



# Milestone 2

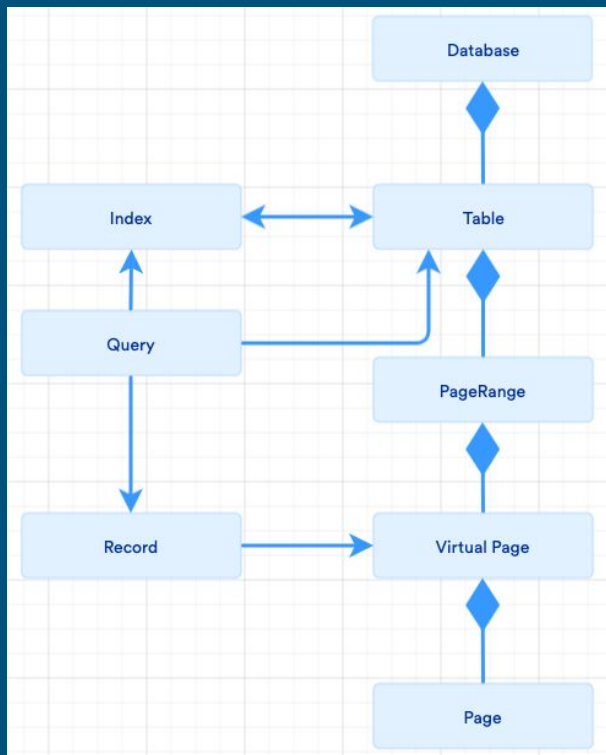


## Merge and Persistent Store

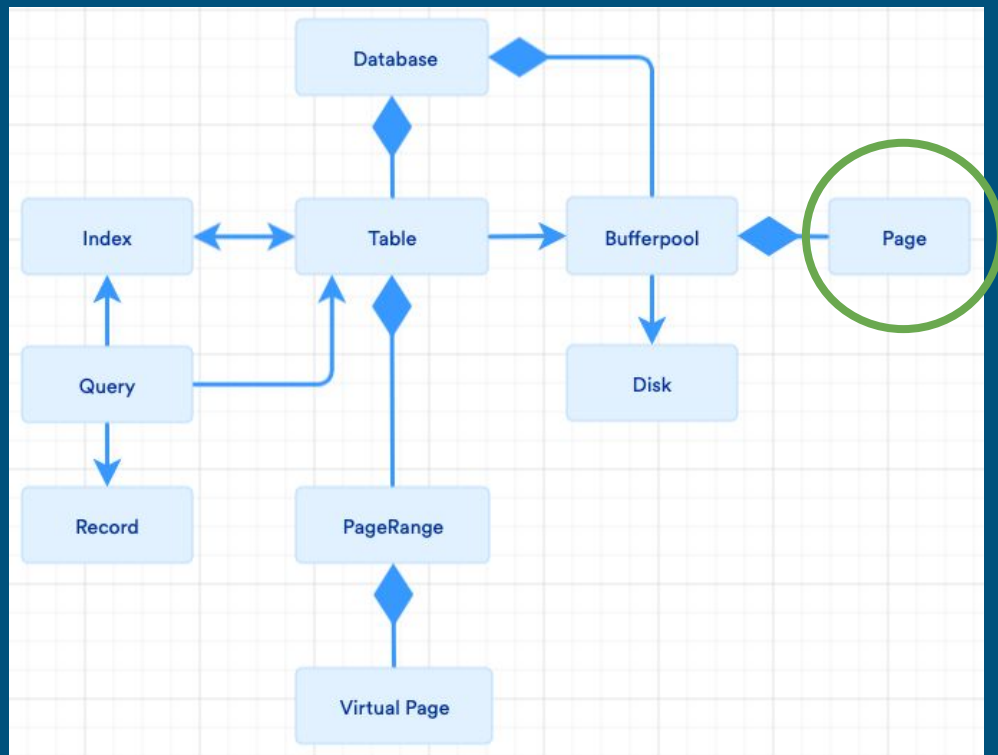
Shivani Parekh, Soumya Durisetti, Alvin Agana, Daniel Silva, Christopher Pires



## Milestone 1



## Milestone 2



# Bufferpool : New Classes

Disk - writes and reads files in a directory specified by database

- One file per physical page
- File name contains identifying information about page

Bufferpool - represents main memory

- Stores Page objects
- Eviction policy based on Clock
- Keeps track of dirty pages and pin counts
- Returns page object to caller (Table class) based on page information
- Gets page from disk if needed

# Bufferpool: Changes to M1

---

## Table class

- Manages pages - creates Pages and requests them from Bufferpool
- When pages are initially created, inserted into Disk
- Calls on bufferpool to get pages from memory or disk, handles pin/unpin

## VirtualPage class

- Stores list of page locations instead of page objects

## Page class

- Has a `page.location` attribute, which is
  - `(table_name, page_range_id, virtual_page_id, page_id/column)`

# Bufferpool: Changes to M1

---

## DB class

- Created a bufferpool instance shared across the entire DB
- `open()`
  - Repopulates tables from disk if any
  - Sets path for bufferpool/disk
  - Repopulates physical page locations using file name structure  
ex: "Grades-0-B\_0-4.txt"
- `close()`
  - All dirty pages in bufferpool are written to disk
  - Table information (page directories, RID directories, Index directory - JSON) are written to disk

# Merge

---

Called in `update()` query for every 100 updates to a base page

- Creates Background-thread that consolidates base pages
- Main-thread continues to handle queries
- Uses Python's Threading Library

# Merge - Implementation

---

## Main-thread

- Default thread which handles regular flow of execution (eg. queries)
- Updates page directory before the start of the next merge process
  - Consolidated base page stored as an attribute of old base page

## Merge-thread (background)

- Consolidates new updates from TID to TPS into a deep copy of the base page
- Once updates are made, background-thread merges with the main-thread

# Index

---

- Implemented as a Hashtable
  - Key is data value
  - Value is list of rids (just one rid for index on primary key)
  - Primary key always has an index
  - Indexes can be created and dropped on different columns
  - One hashtable per column
- Pros
  - Quick access time for single values
  - Easy implementation
- Cons
  - Takes longer to access a range

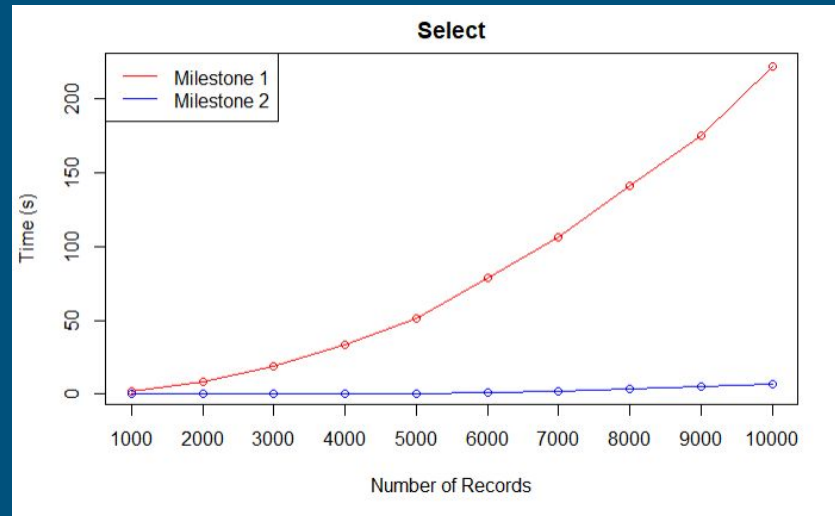



# Coding Quality Improvements from Milestone 1


---

- Unit testing + end-to-end tests: thoroughly test out codebase
- Input validation: makes code more robust and easy to debug
- Standardized function structure and commenting etiquette

# Results



# Results



# Future Improvements

---

- Optimize merge threshold for better performance
- Create more unit tests for easier debugging in the future
- Break large functions into smaller methods for improved readability

# Challenges

---

Initially, we struggled to understand how to implement bufferpool and merge.

- Merge
  - Where and when to call merge function
  - How to have a thread merging records in the background
- Bufferpool
  - How to represent disk/ what to put in the files
  - How to re-populate tables from persistent store

# Takeaways

---

- Better understanding of how persistent store works
  - Didn't realize that not only page data gets stored but also page directory
- When splitting up tasks we should discuss changes/ plan of action together
  - Maybe have a project coordinator role